


Intelligent Systems Programming

Lecture 4: Propositional Logic



Verification and Expert Systems



Today's Program

- **Propositional Logic [12:00-12:50]**
 - Fundamental Concepts in Logic
 - Syntax and Semantics
 - Inference
 - Entailment
 - Logical equivalence
 - Inference rules
 - Formal proofs
- **Efficient Inference Algorithms [13:00-13:50]**
 - Resolution
 - Inference with Horn clauses
 - Efficient SAT checking
 - DPLL
 - WalkSat
 - Phase transition
- **Deloitte Hackathon [13:50-14:00]**
- **Exercises [14:00-16:00]**



Fundamental Concepts of Logic



The Purpose of Logics

- **Logics** are formal languages for representing information such that conclusions can be drawn
- Natural language is too complex and **ambiguous**
 - “John saw the diamond through the window and stole it”
 - Reading 1: John stole the diamond
 - Reading 2: John stole the window
- **Sentences** in logics are assertions about a world that are either true or false

Logic: Syntax and Semantics

- **Syntax** defines the written form of legal **sentences** in the language
- **Semantics** define the truth-value of sentences in a **world**
- **World** is the setting or environment in which you derive the truth of sentences
- E.g., the language of inequalities
 - $x+2 \geq y$ is a sentence; $x^2+y > \{\}$ is not a sentence
 - $x+2 \geq y$ is true in a world where $x = 7, y = 1$
 - $x+2 \geq y$ is false in a world where $x = 0, y = 6$

Entailment and Inference

- **Entailment** means that one thing follows from another:

$$KB \models \alpha$$

\models is a meta symbol, not a part of the syntax!

- Knowledge base **KB** **entails** sentence α if and only if α is true in all worlds where **KB** is true
 - E.g., KB containing *the-apple-is-red* and *the-apple-is-sweet* entails *the-apple-is-sweet*
 - E.g., KB containing “ $y \geq 4$ ”, “ $y \leq 4$ ” entails $y = 4$
- **Inference** is to **decide** whether $KB \models \alpha$

Models

- A **model** is a formal description of a **possible world** used to decide truth-value of sentences

Example:


Possible world = state of nuclear power plant

Model = state *{broken, hot, cold}* of pipe *A, B, and C*

- We say *m* is a **model** of a sentence *α* if *α* is true in *m*
- $M(\alpha)$ is the set of all models of *α*

Models

- $KB \models \alpha$ if and only if $M(KB) \subseteq M(\alpha)$



Inference Algorithms

- $KB \vdash_i \alpha$: sentence α can be derived from KB by procedure i
- **Soundness**: i is sound if $KB \vdash_i \alpha$ implies $KB \models \alpha$
 - Any sentence derived by i from KB is **truth preserving**
- **Completeness**: i is complete if $KB \models \alpha$ implies $KB \vdash_i \alpha$
 - All the sentences entailed by KB can be derived by procedure i
 - That is, the procedure will answer any question whose answer follows from what is known by the KB

Propositional Logic



Syntax

- Atomic sentences

- Proposition symbols P, Q, \dots are sentences
- The two constants *True* and *False* are sentences

- Complex sentences

- If S_1 and S_2 are sentences then so are (in order of precedence)

• $\neg S_1$	<i>negation</i>	\neg <i>not</i>	$\neg Q, Q$ <i>literals</i>
• $(S_1 \wedge S_2)$	<i>conjunction</i>	\wedge <i>and</i>	S_1, S_2 <i>conjuncts</i>
• $(S_1 \vee S_2)$	<i>disjunction</i>	\vee <i>or</i>	S_1, S_2 <i>disjuncts</i>
• $(S_1 \Rightarrow S_2)$	<i>implication</i>	\Rightarrow <i>implies</i>	S_1 <i>premise</i> S_2 <i>conclusion</i>
• $(S_1 \Leftrightarrow S_2)$	<i>biimplication</i>	\Leftrightarrow <i>if-and-only-if</i>	

Semantics (1)

- Each model m assigns **truth value** *true* (1) or *false* (0) to each proposition symbol

E.g.

P	Q	R
<i>false</i>	<i>true</i>	<i>false</i>

- The constants *True* and *False* are always given the expected values

<i>True</i>	<i>False</i>
<i>true</i>	<i>false</i>

- This specifies the truth value of atomic sentences

Semantics (2)

- Rules for evaluating truth with respect to a model m :

$\neg S$	is <i>true</i>	iff	S is <i>false</i>	
$S_1 \wedge S_2$	is <i>true</i>	iff	S_1 is <i>true</i>	and S_2 is <i>true</i>
$S_1 \vee S_2$	is <i>true</i>	iff	S_1 is <i>true</i>	or S_2 is <i>true</i>
$S_1 \Rightarrow S_2$	is <i>true</i>	iff	S_1 is <i>false</i>	or S_1, S_2 are <i>true</i>
$S_1 \Leftrightarrow S_2$	is <i>true</i>	iff	$S_1 \Rightarrow S_2$ is <i>true</i>	and $S_2 \Rightarrow S_1$ is <i>true</i>

- Simple recursive process evaluates an arbitrary sentence, e.g.,

$$\neg P \Rightarrow (Q \wedge R) = \neg \textit{false} \Rightarrow (\textit{true} \wedge \textit{false}) = \textit{true} \Rightarrow \textit{false} = \textit{false}$$

Validity

A sentence is **valid** if it is true in **all** models

e.g., *True*, $A \vee \neg A$, $A \Rightarrow A$, ...

Validity is connected to entailment via the

Deduction Theorem:

$KB \models \alpha$ if and only if $(KB \Rightarrow \alpha)$ is valid

Satisfiability

A sentence is **satisfiable** if it is true in **some** model

e.g., $A \vee B$

A sentence is **unsatisfiable** if it is true in **no** models

e.g., $A \wedge \neg A$

Satisfiability is connected to entailment via the following:

$KB \models \alpha$ if and only if $(KB \wedge \neg \alpha)$ is unsatisfiable

Inference by Enumeration

- Inference: decide whether $KB \models \alpha$

function TT-ENTAILS?(KB, α) **returns** *true* or *false*

inputs: KB , the knowledge base, a sentence in propositional logic
 α , the query, a sentence in propositional logic

$symbols \leftarrow$ a list of the proposition symbols in KB and α

return TT-CHECK-ALL($KB, \alpha, symbols, \{ \}$)

function TT-CHECK-ALL($KB, \alpha, symbols, model$) **returns** *true* or *false*

if EMPTY?($symbols$) **then**

if PL-TRUE?($KB, model$) **then return** PL-TRUE?($\alpha, model$)

else return *true* // when KB is false, always return *true*

else do

$P \leftarrow$ FIRST($symbols$)

$rest \leftarrow$ REST($symbols$)

return (TT-CHECK-ALL($KB, \alpha, rest, model \cup \{P = true\}$)

and

 TT-CHECK-ALL($KB, \alpha, rest, model \cup \{P = false\}$))

Properties of TT-ENTAILS

- DFS enumeration and check of all models
- Sound? yes checks if α is true when KB is true
- Complete? yes checks all models
- For n symbols
 - time complexity is $O(2^n)$
 - space complexity is $O(n)$

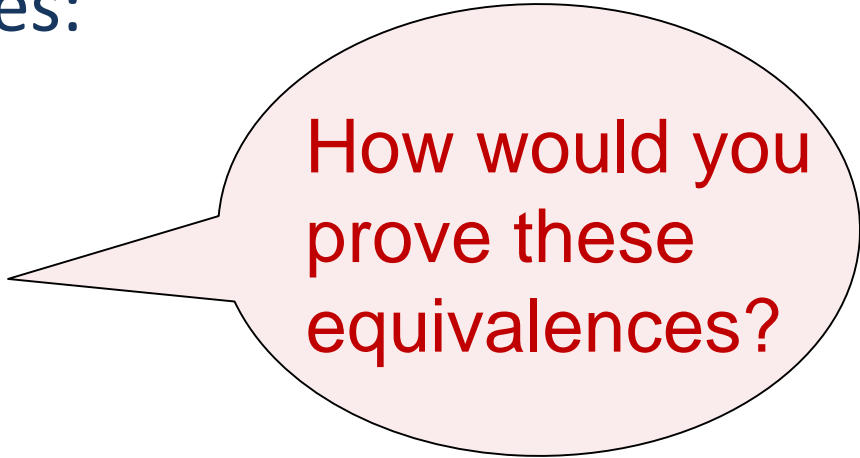
Logical Equivalence

- Two sentences are **logically equivalent** iff they are true in the same models:

$$\alpha \equiv \beta \text{ iff } M(\alpha) = M(\beta) \text{ iff } \alpha \models \beta \text{ and } \beta \models \alpha$$

- Standard equivalences:

- $\alpha \wedge \neg \alpha \equiv \text{False}$
- $\alpha \vee \neg \alpha \equiv \text{True}$
- $\alpha \wedge \text{True} \equiv \alpha$
- $\alpha \vee \text{False} \equiv \alpha$
- $\alpha \wedge \text{False} \equiv \text{False}$
- $\alpha \vee \text{True} \equiv \text{True}$
- $\alpha \wedge \alpha \equiv \alpha$
- $\alpha \vee \alpha \equiv \alpha$



How would you prove these equivalences?

More Standard Equivalences

$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha)$ commutativity of \wedge

$(\alpha \vee \beta) \equiv (\beta \vee \alpha)$ commutativity of \vee

$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma))$ associativity of \wedge

$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma))$ associativity of \vee

$\neg(\neg\alpha) \equiv \alpha$ double-negation elimination

$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha)$ contraposition

$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta)$ implication elimination

$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha))$ biconditional elimination

$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta)$ de Morgan

$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta)$ de Morgan

$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma))$ distributivity of \wedge over \vee

$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma))$ distributivity of \vee over \wedge

Sentence Simplification

- Determine logical equivalence between sentences using standard rules

Example

$$\begin{aligned} & (a \vee (b \Rightarrow a)) \\ \equiv & (a \vee (\neg b \vee a)) \quad \text{impl. elim.} \\ \equiv & \dots \\ \equiv & b \Rightarrow a \end{aligned}$$

- Formal proof!

Inference Rules

numerator (premises) \vdash denominator (conclusion)

$$\frac{\alpha \Rightarrow \beta, \alpha}{\beta}$$

Modus Ponens

$$\frac{\alpha \wedge \beta}{\alpha}$$


And-Elimination

$$\frac{\alpha \Rightarrow \beta}{\neg \alpha \vee \beta} \quad \frac{\neg \alpha \vee \beta}{\alpha \Rightarrow \beta}$$

All Equivalence rules

Inference Proof

- Search for inference rules to chain “goal” with axioms



Only complete approach if set of inference rules is complete!

Efficient Inference Algorithms



Conjunctive Normal Form (CNF)

Definition

- A **literal** is a symbol or a negated symbol
- A **clause** is a disjunction of literals
- **CNF** is a conjunction of clauses

E.g., $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$

Translate Arbitrary Sentence to CNF

1. Eliminate \Leftrightarrow , replacing $\alpha \Leftrightarrow \beta$ with $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$
2. Eliminate \Rightarrow , replacing $\alpha \Rightarrow \beta$ with $\neg \alpha \vee \beta$
3. Move \neg inwards using de Morgan's rules and double-negation
4. Apply distribution law (\wedge over \vee) and flatten
5. Eliminate symbol duplicates in clauses (**factoring**)

Resolution: A Complete Inference Rule

$$\frac{\ell_1 \vee \dots \vee \ell_k, \quad m_1 \vee \dots \vee m_n}{\ell_1 \vee \dots \vee \ell_{i-1} \vee \ell_{i+1} \vee \dots \vee \ell_k \vee m_1 \vee \dots \vee m_{j-1} \vee m_{j+1} \vee \dots \vee m_n}$$

where ℓ_i and m_j are complementary literals ($\ell_i \equiv \neg m_j$).

$$\text{E.g., } \frac{P \vee Q, \quad \neg Q \vee R}{P \vee R}$$

Resolution Algorithm

- Remember: $KB \models \alpha$ iff $KB \wedge \neg \alpha$ is unsatisfiable
- Keep doing resolution on clauses until
 - fixpoint reached with no empty clause $()$
return *false*
 - $()$ (= False) derived
return *true*

$$\frac{Q, \neg Q}{()}$$

Resolution Algorithm

function PL-RESOLUTION(KB, α) **returns** *true* or *false*

inputs: KB , the knowledge base, a sentence in propositional logic
 α , the query, a sentence in propositional logic

$clauses \leftarrow$ the set of clauses in the CNF representation of $KB \wedge \neg\alpha$

$new \leftarrow \{ \}$

loop do

for each pair of clauses C_i, C_j **in** $clauses$ **do**

$resolvents \leftarrow$ PL-RESOLVE(C_i, C_j)

if $resolvents$ contains the empty clause **then return** *true*

$new \leftarrow new \cup resolvents$

if $new \subseteq clauses$ **then return** *false*

$clauses \leftarrow clauses \cup new$

Properties of the Resolution Algorithm

- **Sound**, yes
- **Complexity**
 - Time, Space: $O(2^n)$, due to blow-up of number of clauses
- **Complete**, yes

Ground resolution theorem

if a CNF with set of clauses S is unsatisfiable, then the **resolution closure** of those clauses, $RC(S)$, contains the empty clause $()$.

Proof of ground resolution theorem

We must show: if S is unsatisfiable then $()$ in $RC(S)$

Prove log. equiv. sentence: if $()$ not in $RC(S)$ then S satisfiable

Satisfying assignment to propositions P_1 to P_k

For $i = 1$ to k

- $P_i = \text{False}$ if there is a clause $(\text{False} \vee \dots \vee \text{False} \vee \neg P_i)$
- $P_i = \text{True}$ otherwise

Why is this assignment satisfying?

Horn Form Clauses

- **Horn clause**: clause with *at most* one positive literal
- **Definit clause**: *exactly* one positive literal

$$\begin{array}{lcl} P_1 \wedge P_2 \wedge \dots \wedge P_n \Rightarrow C & \text{---} & \text{head} \\ \Rightarrow F & \text{---} & \text{body} \\ & \text{---} & \text{fact} \end{array}$$

- **Modus Ponens** (for Definit clause)

$$\frac{\alpha_1, \dots, \alpha_n, \quad \alpha_1 \wedge \dots \wedge \alpha_n \Rightarrow \beta}{\beta}$$

- **Inference possible in linear time!**

Forward chaining inference

KB:

$P \Rightarrow Q$

$L \wedge M \Rightarrow P$

$B \wedge L \Rightarrow M$


$A \wedge P \Rightarrow L$

$A \wedge B \Rightarrow L$


A

B


$KB \models Q ?$




Forward chaining example




Forward chaining example




Forward chaining example




Forward chaining example




Forward chaining example



Forward chaining example




Forward chaining example



Forward chaining example

$KB \models Q$

Yes!



Properties of Forward Chaining

- **Sound**, yes since Modus Ponens is sound
- **Complete**, yes
- **Space and time**: $O(n)$, where n is the total number of clause literals

Efficient SAT-Checking

Complete backtracking search algorithms

- **DPLL algorithm** (Davis, Putnam, Logemann, Loveland)

Incomplete local search algorithms

- **WalkSAT** algorithm

The DPLL Algorithm

Determine if a CNF sentence is satisfiable

Improvements over **truth table enumeration**:

1. Early termination

A clause is true if any literal is true

A sentence is false if any clause is false

2. Pure symbol heuristic

Pure symbol: always appears with the same "sign" in all clauses

e.g., In the three clauses $(A \vee \neg B)$, $(\neg B \vee \neg C)$, $(C \vee A)$, A and B are pure, C is impure

Make a pure symbol literal true

3. Unit propagation

Unit clause: only one literal in the clause

The only literal in a unit clause must be true

e.g., $(False \vee \neg B)$: B must be false

The DPLL Algorithm

function DPLL-SATISFIABLE?(*s*) **returns** *true* or *false*

inputs: *s*, a sentence in propositional logic

clauses \leftarrow the set of clauses in the CNF representation of *s*

symbols \leftarrow a list of the proposition symbols in *s*

return DPLL(*clauses*, *symbols*, { })

function DPLL(*clauses*, *symbols*, *model*) **returns** *true* or *false*

if every clause in *clauses* is true in *model* **then return** *true*

if some clause in *clauses* is false in *model* **then return** *false*

P, *value* \leftarrow FIND-PURE-SYMBOL(*symbols*, *clauses*, *model*)

if *P* is non-null **then return** DPLL(*clauses*, *symbols* - *P*, *model* \cup { *P*=*value* })

P, *value* \leftarrow FIND-UNIT-CLAUSE(*clauses*, *model*)

if *P* is non-null **then return** DPLL(*clauses*, *symbols* - *P*, *model* \cup { *P*=*value* })

P \leftarrow FIRST(*symbols*); *rest* \leftarrow REST(*symbols*)

return DPLL(*clauses*, *rest*, *model* \cup { *P*=*true* }) **or**

DPLL(*clauses*, *rest*, *model* \cup { *P*=*false* })



The WalkSAT algorithm

Determine if a CNF sentence is satisfiable

Algorithm

- Start with a random complete assignment
- In each iteration:
 - Pick random false clause
 - With probability p
 - flip random literal in clause
 - Else
 - flip literal that makes most clauses true
- Stop when reaching given max number of iterations



Phase Transition

- Consider random 3-CNF sentences. e.g.,


$$(\neg D \vee \neg B \vee C) \wedge (B \vee \neg A \vee \neg C) \wedge (\neg C \vee \neg B \vee E) \\ \wedge (E \vee \neg D \vee B) \wedge (B \vee E \vee \neg C)$$

m = number of clauses

n = number of symbols

- Hard problems seem to cluster near $m/n = 4.3$
(phase transition)

Phase Transition



Phase Transition

